

CSE 323: Operating System Design


Locking

Salman Shamil



North South University (NSU)
Fall 2025

Original slides by Mathias Payer and Sanidhya Kashyap [EPFL]



- Abstraction: locks to protect shared data structures
- Mechanism: interrupt-based locks
- Mechanism: atomic hardware locks
- Busy waiting (spin locks) versus wait queues

This slide deck covers chapters 28 in OSTEP.

Race Conditions

- Concurrent execution leads to race conditions
 - Access to shared data must be mediated
- **Critical section:** part of code that accesses shared data
- **Mutual exclusion:** only one process is allowed to execute critical section at any point in time
- **Atomicity:** critical section executes as an uninterruptible block

A **mechanism** to achieve atomicity is through locking.

Locks: Basic Idea

- Lock variable protects critical section
- All threads competing for *critical section* share a lock
- Only one thread succeeds at acquiring the lock (at a time)
- Other threads must wait until lock is released

```
lock_t mutex;  
...  
lock(&mutex);  
cnt = cnt + 1;  
unlock(&mutex);
```

- Requirements: mutual exclusion, fairness, and performance
 - **Mutual exclusion:** only one thread in critical section
 - **Fairness:** all threads should eventually get the lock
 - **Performance:** low overhead for acquiring/releasing lock
- Lock implementation requires hardware support
 - ... and OS support for performance

Lock Operations

- `void lock(lock_t *lck):` acquires the lock, current thread owns the lock when function returns
- `void unlock(lock_t *lck):` releases the lock

Lock Operations

- `void lock(lock_t *lck)`: acquires the lock, current thread owns the lock when function returns
- `void unlock(lock_t *lck)`: releases the lock

Note that we assume that the application *correctly* uses locks for *each* access to the critical section.

Interrupting Locks

- Turn off interrupts when executing critical sections
 - Neither hardware nor timer can interrupt execution
 - Prevent scheduler from switching to another thread
 - Code between interrupts executes atomically

```
void acquire(lock_t *l) {  
    disable_interrupts();  
}
```

```
void release(lock_t *l) {  
    enable_interrupts();  
}
```

Interrupting Locks (Disadvantages)

- No support for locking multiple locks
- Only works on uniprocessors (no support for locking across cores in multicore system)
- Process may keep lock for arbitrary length
- Hardware interrupts may get lost (hardware only stores information that interrupt X happened, not how many times it happened)

Interrupting Locks (Perspective)

- Interrupt-based locks are extremely simple
- Work well for low-complexity code

Interrupting Locks (Perspective)

- Interrupt-based locks are extremely simple
- Work well for low-complexity code
- Implementing locks through interrupts is great for MCUs

(Faulty) Spin Lock

- Use a shared variable to synchronize access to critical section

```
bool lock1 = false;
```

```
void acquire(bool *lock) {  
    while (*lock); /* spin until we grab the lock */  
    *lock = true;  
}
```

```
void release(bool *lock) {  
    *lock = false  
}
```

(Faulty) Spin Lock

- Use a shared variable to synchronize access to critical section

```
bool lock1 = false;
```

```
void acquire(bool *lock) {  
    while (*lock); /* spin until we grab the lock */  
    *lock = true;  
}
```

```
void release(bool *lock) {  
    *lock = false  
}
```

Bug: both threads can grab the lock if thread is preempted before setting the lock but after the while loop completes.

Locking requires an atomic *test-and-set* instruction.

```
int TestAndSet(int *addr, int val) {  
    int old = *addr;  
    *addr = val;  
    return old;  
}
```

This pseudocode in **c** demonstrates the basic idea of an atomic exchange instruction (xchg on x86 or ldstub on SPARC).

Test-and-set Spin Lock

```
int lock1; // 0 -> lock is available, 1 -> lock is held

void acquire(int *lock) {
    while (TestAndSet(lock, 1) == 1); /* spin */
}

void release(int *lock) {
    *lock = 0;
}

acquire(&lock1);
critical_section();
release(&lock1);
```

This time we guarantee that the thread that changes lock from 0 to 1 gets to execute its critical section.

Compare-and-swap Spin Lock

```
int CompareAndSwap(int *ptr, int expt, int new) {  
    int actual = *ptr;  
    if (actual == expt) {  
        *ptr = new;  
    }  
    return actual;  
}
```

Compare-and-swap Spin Lock

```
int CompareAndSwap(int *ptr, int expt, int new) {  
    int actual = *ptr;  
    if (actual == expt) {  
        *ptr = new;  
    }  
    return actual;  
}
```

- Returns the actual value (before the potential update), indicating whether it succeeded or not.
- More powerful than test-and-set [blind vs conditional update]

Compare-and-swap Spin Lock

```
int CompareAndSwap(int *ptr, int expt, int new) {  
    int actual = *ptr;  
    if (actual == expt) {  
        *ptr = new;  
    }  
    return actual;  
}
```

- Returns the actual value (before the potential update), indicating whether it succeeded or not.
- More powerful than test-and-set [blind vs conditional update]

```
void acquire_cas(int *lock) {  
    while (CompareAndSwap(lock, 0, 1) == 1); /* spin */  
}
```

Ticket Lock with Fetch-And-Add

- Neither test-and-set nor compare-and-swap guarantees progress.
 - A thread may spin forever.

Ticket Lock with Fetch-And-Add

- Neither test-and-set nor compare-and-swap guarantees progress.
 - A thread may spin forever.
- Another hardware primitive Fetch-And-Add

```
int FetchAndAdd(int *ptr) {  
    int old = *ptr;  
    *ptr = old + 1;  
    return old;  
}
```

Ticket Lock with Fetch-And-Add

- Neither test-and-set nor compare-and-swap guarantees progress.
 - A thread may spin forever.
- Another hardware primitive Fetch-And-Add

```
int FetchAndAdd(int *ptr) {  
    int old = *ptr;  
    *ptr = old + 1;  
    return old;  
}
```

- Fetch-And-Add can be used to build Ticket Lock, where a thread once queued, will eventually acquire the lock.

```
typedef struct __lock_t {  
    int ticket;  
    int turn;  
} lock_t;
```

Ticket Lock with Fetch-And-Add

```
void lock_init(lock_t *lock) {
    lock->ticket = 0;
    lock->turn = 0;
}

void lock(lock_t *lock) {
    // get my ticket
    int myturn = FetchAndAdd(&lock->ticket);
    while (lock->turn != myturn) {
        ; // spin until it's my turn
    }
}

void unlock(lock_t *lock) {
    // next ticket goes
    lock->turn = lock->turn + 1;
}
```

Spin Lock: Reduce Spinning

- A simple way to reduce the cost of spinning is to `yield()` whenever lock acquisition fails
 - This is no longer a “strict” spin lock as we give up control to the scheduler every loop iteration

```
void acquire(bool *lck) {  
    while (TestAndSet(1, 1) == 1) {  
        yield();  
    }  
}
```


A Better Way: Queue Lock

- Idea: instead of spinning, put threads on a queue
- Wake up thread(s) when lock is released
 - Wake up all threads to have them race for the lock
 - Selectively wake one thread up for fairness
- OS Support: `park()` and `unpark(threadID)`