

CSE 323: Operating System Design

Condition Variable & Semaphore

Salman Shamil



North South University (NSU)
Fall 2025

Original slides by Mathias Payer and Sanidhya Kashyap [EPFL]

Salman Shamil

CSE 323: Operating System Design

1 / 28

Lecture Topics

- Condition Variables
- Producer-Consumer Problem
- Semaphores
- Signaling through condition Variables and Semaphores
- Concurrency Bugs

This slide deck covers chapters 30 and 31 in OSTEP.

Salman Shamil

CSE 323: Operating System Design

2 / 28

Condition Variables (CV)

In concurrent programming, a common scenario is one thread waiting for another thread to complete an action.

```
1 bool done = false;
2
3 /* called in the child to signal termination */
4 void thr_exit() {
5     done = true;
6 }
7 /* called in the parent to wait for a child thread */
8 void thr_join() {
9     while (!done);
10 }
```

Condition Variables (CV)

- Locks enable mutual exclusion of a shared region.
 - Unfortunately they are oblivious to ordering
- Waiting and signaling (i.e., T2 waits until T1 completes a given task) could be implemented by spinning until the value changes
- But spinning is incredibly *inefficient*
- New synchronization primitive: **condition variables**

Salman Shamil

CSE 323: Operating System Design

3 / 28

Salman Shamil

CSE 323: Operating System Design

4 / 28

Condition Variables (CV)

- A CV allows:
 - A thread to wait for a condition
 - Another thread signals the waiting thread
- Implement CV using queues
- API: wait, signal or broadcast
 - wait: wait until a condition is satisfied
 - signal: wake up one waiting thread
 - broadcast: wake up all waiting threads
- On Linux, pthreads provides CV implementation

Signal parent that child has exited

```
1 bool done = false;
2 pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
3 pthread_cond_t c = PTHREAD_COND_INITIALIZER;
4 /* called in the child to signal termination */
5 void thr_exit() {
6     pthread_mutex_lock(&m);
7     done = true;
8     pthread_cond_signal(&c);
9     pthread_mutex_unlock(&m);
10 }
11 /* called in the parent to wait for a child thread */
12 void thr_join() {
13     pthread_mutex_lock(&m);
14     while (!done)
15         pthread_cond_wait(&c, &m);
16     pthread_mutex_unlock(&m);
17 }
```

Signal parent that child has exited (2)

- pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m)
 - Assume mutex m is held; *atomically* unlock mutex when waiting, retake it when waking up
- Question: Why do we need to check a condition before sleeping?
- Thread may have already exited, i.e., no need to wait
 - Principle: Check the condition before sleeping
- Question: Why can't we use if when waiting?
- There can be multiple threads to wait on the same CV. **Race Condition!**
 - Principle: while instead of if when waiting (more on this later...)

Signal parent that child has exited (3)

- Question: Why do we need to protect done with mutex m?
- Mutex m allows one thread to access done for protecting against missed updates
 - Parent reads done == false but is interrupted
 - Child sets done = true and signals but no one is waiting
 - Parent continues and goes to sleep (forever)
- Lock is therefore required for wait/signal synchronization

Producer/Consumer Problem

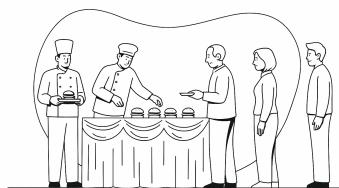


Figure 1: Producer-Consumer/Bounded Buffer Problem

- Producer/consumer is a common programming pattern
- For example: map (producers) / reduce (consumer)
- For example: a concurrent database (consumers) handling parallel requests from clients (producers)
 - Clients produce new requests (encoded in a queue)
 - Handlers consume these requests (popping from the queue)

Producer/Consumer with Bounded Buffer

- One or more producers create items, store them in buffer
- One or more consumers process items from buffer
- Need synchronization for buffer
 - Want concurrent production and consumption
 - Use as many cores as available
 - Minimize access time to shared data structure
- Strategy: use CV to synchronize
 - Make producers wait if buffer is full
 - Make consumers wait if buffer is empty (nothing to consume)

Solving Producer/Consumer Problem

• Setup:

- Buffer holds a single item
- One producer and one consumer

```
int buffer;
int count = 0; // initially empty

void put(int value) {
    assert(count == 0);
    count = 1;
    buffer = value;
}

int get() {
    assert(count == 1);
    count = 0;
    return buffer;
}

void *producer(void *arg) {
    int i;
    int loops = (int) arg;
    for (i = 0; i < loops; i++) {
        put(i);
    }
}

void *consumer(void *arg) {
    while (1) {
        int tmp = get();
        printf("%d\n", tmp);
    }
}
```

• Problems with this solution

- Critical sections in `put()` and `get()`. **Use locks...**
- Producer-Consumer dependency for fetching. **Needs CV!**

Solving Producer/Consumer Problem

```
cond_t cond;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    int loops = (int) arg;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        if (count == 0)
            Pthread_cond_wait(&cond, &mutex);
        Pthread_mutex_lock(&mutex);
        if (count == 1)
            Pthread_cond_wait(&cond, &mutex);
        put(i);
        Pthread_cond_signal(&cond);
        Pthread_mutex_unlock(&mutex);
    }
}

void *consumer(void *arg) {
    int i;
    int loops = (int) arg;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        if (count == 0)
            Pthread_cond_wait(&cond, &mutex);
        int tmp = get();
        Pthread_cond_signal(&cond);
        Pthread_mutex_unlock(&mutex);
        printf("%d\n", tmp);
    }
}
```

Does it work?

- Fine for single producer and single consumer.
- Change the setup to accommodate multiple producers and/or multiple consumers. How about now?

Solving Producer/Consumer Problem (2)

- **Setup:**
 - Buffer holds a single item
 - Multiple producers and/or multiple consumers
- Previous code doesn't work.

No recheck after waking up.

- Consider a consumer thread (C1) is waiting for an item
- What if a second consumer thread (C2) sneaks in just after an item is produced? ... skipping the `wait()` call.
- Producer's `signal()` wakes C1 up, but C2 already fetched the item!
- Solution: Use `while` instead of `if` to recheck upon waking up.

Producers and consumers both waiting on the same CV.

- Two consumers C1 and C2 runs and sleeps by calling `wait()`.
- Producer runs and `signal()` wakes up C1 (or C2).
- After consuming the item C1 can wake up producer again.
- But what if C1's `signal()` wakes up C2 instead?
- Solution: Use separate conditions for directed signaling.

Solving Producer/Consumer Problem (2)

```
cond_t empty, full;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    int loops = (int) arg;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        while (count == 0)
            Pthread_cond_wait(&empty, &mutex);
        count++;
        Pthread_mutex_unlock(&mutex);
        Pthread_cond_signal(&full);
    }
}

void *consumer(void *arg) {
    int i;
    int loops = (int) arg;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        while (count == 1)
            Pthread_cond_wait(&full, &mutex);
        int tmp = get();
        Pthread_mutex_unlock(&mutex);
        Pthread_cond_signal(&empty);
        printf("%d\n", tmp);
    }
}
```

Producer/Consumer Buffer with Multiple Slots

```
int buffer[MAX];
int fill_ptr = 0;
int use_ptr = 0;
int count = 0;

void put(int value) {
    buffer[fill_ptr] = value;
    fill_ptr = (fill_ptr + 1) % MAX;
    count++;
}

int get() {
    int tmp = buffer[use_ptr];
    use_ptr = (use_ptr + 1) % MAX;
    count--;
    return tmp;
}

cond_t empty, full;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        while (count == MAX)
            Pthread_cond_wait(&empty, &mutex);
        put(i);
        Pthread_cond_signal(&full);
        Pthread_mutex_unlock(&mutex);
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        while (count == 0)
            Pthread_cond_wait(&full, &mutex);
        int tmp = get();
        Pthread_mutex_unlock(&mutex);
        Pthread_cond_signal(&empty);
        printf("%d\n", tmp);
    }
}
```

Semaphore

- A semaphore extends a CV with an integer as internal state
- `int sem_init(sem_t *sem, unsigned int value):` creates a new semaphore with value slots
- `int sem_wait(sem_t *sem):` waits until the semaphore has at least one slot, decrements the number of slots
- `int sem_post(sem_t *sem):` increments the semaphore (and wakes one waiting thread)
- `int sem_destroy(sem_t *sem):` destroys the semaphore and releases any waiting threads

Producer/Consumer: Use Semaphores!

```
sem_t csem, psem;

/* BUFSIZE items are available for producer to create */
sem_init(&psem, 0, BUFSIZE);

/* 0 items are available for consumer */
sem_init(&csem, 0, 0);
```

Producer: Semaphores

```
1 void put(unsigned int val) {
2     /* we wait until there is buffer space available */
3     sem_wait(&psem);
4
5     /* store element in buffer */
6     buffer[ppos] = val;
7     ppos = (ppos + 1) % BUFSIZE;
8
9     /* notify consumer that data is available */
10    sem_post(&csem);
11 }
```

Consumer: Semaphores

```
1 unsigned int get() {
2     /* wait until data is produced */
3     sem_wait(&csem);
4
5     /* consumer entry */
6     unsigned long val = buffer[cpos];
7     cpos = (cpos + 1) % BUFSIZE;
8
9     /* notify producer that a space has freed up */
10    sem_post(&psem);
11    return val;
12 }
```

Producer/Consumer: Remaining Issues?

- We now synchronize between consumers and producers
 - Producer waits until buffer space is available
 - Consumer waits until data is ready
- How would you handle multiple producers/consumers?
 - Currently no synchronization between producers (or consumers)

Multiple Producers: Use Locking!

```
/* mutex handling mutual exclusive access to ppos */
1 pthread_mutex_t pmutex = PTHREAD_MUTEX_INITIALIZER;
2
3 void put(unsigned int val) {
4     unsigned int mypos;
5     /* we wait until there is buffer space available */
6     sem_wait(&psem);
7     /* ppos is shared between all producers */
8     pthread_mutex_lock(&pmutex);
9     mypos = ppos;
10    ppos = (ppos + 1) % BUFSIZE;
11    /* store information in buffer */
12    buffer[mypos] = val;
13    pthread_mutex_unlock(&pmutex);
14    sem_post(&csem);
15 }
```

Semaphores/Spin Locks/CVs are interchangeable

- Each is implementable through a combination of the others
- Depending on the use-case one is faster than the other
 - How often is the critical section executed?
 - How many threads compete for a critical section?
 - How long is the lock taken?

Implementing a Mutex with a Semaphore

```
1 sem_t sem;
2 sem_init(&sem, 1);
3
4 sem_wait(&sem);
5 ... // critical section
6 sem_post(&sem);
```

Implementing a Semaphore with CV/Locks

```
1 typedef struct {
2     int value;           // sem value
3     pthread_mutex_t lock; // access to sem
4     pthread_cond_t cond; // wait queue
5 } sem_t;
6
7 void sem_init(sem_t *s, int val) {
8     s->value = val;
9     pthread_mutex_init(&(s->lock), NULL);
10    pthread_cond_init(&(s->cond), NULL);
11 }
```

Implementing a Semaphore with CV/Locks

```
1 void sem_wait(sem_t *s) {
2     pthread_mutex_lock(&(s->lock));
3     while (s->value <= 0)
4         pthread_cond_wait(&(s->cond), &(s->lock));
5     s->value--;
6     pthread_mutex_unlock(&(s->lock));
7 }
8
9 void sem_post(sem_t *s) {
10    pthread_mutex_lock(&(s->lock));
11    s->value++;
12    pthread_cond_signal(&(s->cond));
13    pthread_mutex_unlock(&(s->lock));
14 }
```

Reader/Writer Locks

- A single (exclusive) writer, multiple (N) concurrent readers
- Implement using two semaphores: lock for the data structure, wlock for the writer
 - Both semaphores initialized with (1)
 - Writer only waits/posts on wlock when acquiring/releasing
 - Reader waits on lock, increments/decrements reader count
 - If number of readers==0, must wait/post on wlock

Reader/Writer Locks

```
1 void rwlock_acquire_readlock(rwlock_t *rw) {
2     sem_wait(&rw->lock);
3     rw->readers++;
4     if (rw->readers == 1)
5         sem_wait(&rw->wlock); // first r, also grab wlock
6     sem_post(&rw->lock);
7 }
8
9 void rwlock_release_readlock(rwlock_t *rw) {
10    sem_wait(&rw->lock);
11    rw->readers--;
13    if (rw->readers == 0)
14        sem_post(&rw->wlock); // last r, also release wlock
15    sem_post(&rw->lock);
16 }
```

- Full Code in Textbook (Page 12)

Summary

- Spin lock, CV, and semaphore synchronize multiple threads
 - Spin lock: atomic access, no ordering, spinning
 - Condition variable: atomic access, queue, OS primitive
 - Semaphore: shared access to critical section with (int) state
- All three primitives are equally powerful
 - Each primitive can be used to implement both other primitives
 - Performance may differ!