

CSE 323: Operating System Design

Thread

Salman Shamil



North South University (NSU)
Fall 2025

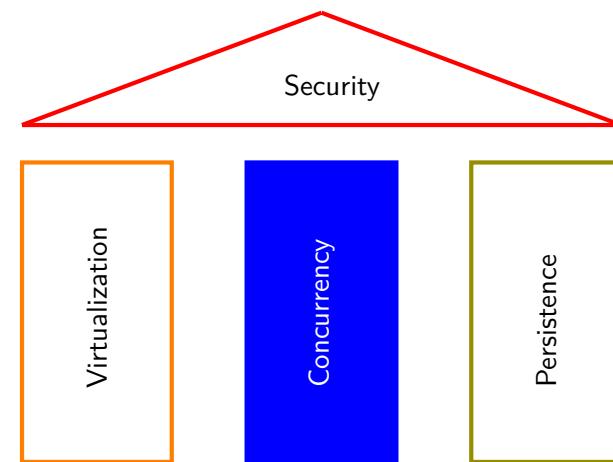
Original slides by Mathias Payer and Sanidhya Kashyap [EPFL]

Salman Shamil

CSE 323: Operating System Design

1 / 10

Concurrency



Salman Shamil

CSE 323: Operating System Design

2 / 10

Lecture Topics

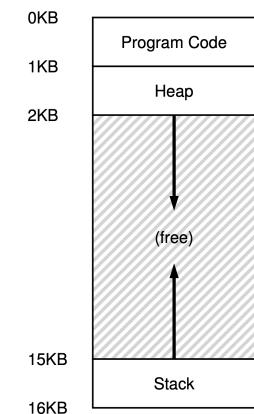
- Thread abstraction
- Multi-threading challenges
- Key concurrency terms and definitions

This slide deck covers chapters 26 and 27 in OSTEP.

[**Credits:** Portions of the content are adapted from slides based on the OSTEP book by Prof. Youjip Won (Hanyang University) and Prof. Mythili Vutukuru (IIT Bombay), with thanks.]

Threads: Execution Contexts

- Threads are independent execution contexts
 - similar to processes
 - EXCEPT they share the same address space
- We only had one thread in a process so far
 - single-threaded program
 - one Program Counter (PC)
 - one Stack Pointer (SP)



Salman Shamil

CSE 323: Operating System Design

3 / 10

Salman Shamil

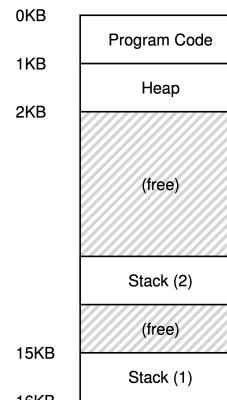
CSE 323: Operating System Design

4 / 10

Multi-threaded Process

- What happens if we want multiple threads in parallel?

- shared address space but separate execution stream
- is that possible with a shared stack or PC?
- each thread has separate stack and PC
 - leading to independent function calls
 - able to execute different parts
- code and heap segments are still shared



- **user-level threads**: scheduled by thread library in user space
- **kernel-level threads**: scheduled directly by the OS

Threads & Concurrency

Concurrency vs Parallelism

- **Concurrency**: multiple processes/threads making progress during the same time period
 - Possibly on a single core by interleaving executions
 - Better CPU utilization (e.g., when one thread is blocked on I/O, another runs)
- **Parallelism**: running multiple processes in parallel over multiple CPU cores
 - A single process can achieve parallelism with multiple threads

How do they communicate?

- Processes need complicated Inter-Process Communication
- Extra memory footprint for IPC
- Threads can do it by simply using global variables (shared)
- **Question:** When to use threads vs processes?

Creating Threads

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

#include "common.h"
#include "common_threads.h"

void *mythread(void *arg) {
    printf("%s\n", (char *) arg);
    return NULL;
}

int main(int argc, char *argv[]) {
    if (argc != 1) {
        fprintf(stderr, "usage: main\n");
        exit(1);
    }

    pthread_t p1, p2;
    printf("main: begin\n");
    Pthread_create(&p1, NULL, mythread, "A");
    Pthread_create(&p2, NULL, mythread, "B");
    // join waits for the threads to finish
    Pthread_join(p1, NULL);
    Pthread_join(p2, NULL);
    printf("main: end\n");
    return 0;
}
```

Shared data is useful but not so simple!

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include "common.h"
#include "common_threads.h"

// shared global variables
int max;
volatile int counter = 0;
// ^ no caching on register

void *mythread(void *arg) {
    char *letter = arg;
    int i; // on stack
    // (private per thread)
    printf("%s: begin \n"
           "[addr of i: %p]\n",
           letter, &i);
    for (i = 0; i < max; i++) {
        counter = counter + 1;
        // shared: only one
    }
    printf("%s: done \n"
           "[counter: %d] \n"
           "[should: %d]\n",
           letter, counter, max*2);
    return 0;
}

int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, \
                "usage: main-first <loopcount>\n");
        exit(1);
    }
    max = atoi(argv[1]);

    pthread_t p1, p2;
    printf("main: begin \n"
           "[counter = %d]\n", counter);
    Pthread_create(&p1, NULL, mythread, "A");
    Pthread_create(&p2, NULL, mythread, "B");
    // join waits for the threads to finish
    Pthread_join(p1, NULL);
    Pthread_join(p2, NULL);
    printf("main: done \n"
           "[counter: %d] \n"
           "[should: %d]\n",
           counter, max*2);
    return 0;
}

Will the final count always be 2 × max?
```

Uncontrolled Scheduling

- assembly instructions for counter = counter + 1 (in x86)

```
100  mov 0x8049a1c, %eax  
105  add $0x1, %eax  
108  mov %eax, 0x8049a1c
```

[CriticalSection] consider a context switch after 'add'.

OS	Thread 1	Thread 2	(after instruction)		
			PC	eax	counter
	before critical section		100	0	50
	mov 8049a1c,%eax		105	50	50
	add \$0x1,%eax		108	51	50
interrupt			100	0	50
save T1			105	50	50
restore T2			108	51	50
	mov 8049a1c,%eax		113	51	51
interrupt			100	0	50
save T2			105	50	50
restore T1			108	51	51
	mov %eax,8049a1c		113	51	51

Concurrency Terms

Race Condition

Concurrent execution of threads leading to different results depending on the order of execution. Such programs are *indeterminate*, producing different outputs across runs.

Critical Section

Portion of code resulting in a race condition, usually by accessing a *shared resource* (e.g., a variable or data structure).

Mutual Exclusion

Guarantees a single thread executes a critical section at a time, preventing race conditions. [Atomicity]

Next: We need to design synchronization primitives for **mutex**.